

# etl

July 6, 2021

```
[1]: import numpy as np
import pandas as pd
```

## 1 Data Preprocessing for articles file

```
[2]: articles = pd.read_json('raw_data/articles/articles_07_04_2021.json')
```

```
[3]: articles.sample(5)
```

```
[3]:          author          linkOfAuthorProfile \
6190      Mohan Gupta  https://towardsdatascience.com/@mohangupta13?s...
31346    Charmaine Chui  https://towardsdatascience.com/@geek-cc?source...
36150      German Osin  https://towardsdatascience.com/@gosin?source=c...
8613  Federico Riveroll  https://towardsdatascience.com/@federicorivero...
20745      Shai Ardazi  https://towardsdatascience.com/@shaiardazi?sou...
```

```
          articleTitle \
6190  A Review of Named Entity Recognition (NER) Usi...
31346  Using Turf.js to Geocode coordinates with cust...
36150  Feature Store as a Foundation for Machine Lear...
8613  Outstanding results predicting Apple Stock app...
20745  Web scraping with Python-A to Z
```

```
          articleLink  postingTime \
6190  https://towardsdatascience.com/a-review-of-nam...  Jul 9, 2018
31346  https://towardsdatascience.com/using-turf-js-t...  Jun 22
36150  https://towardsdatascience.com/feature-store-a...  Dec 10, 2020
8613  https://towardsdatascience.com/making-a-contin...  Feb 13, 2020
20745  https://towardsdatascience.com/web-scraping-wi...  Feb 7, 2019
```

```
          minToRead  recommendations  responses
6190  11 min read          724  11 responses
31346   3 min read           1         None
36150  12 min read          483   2 responses
8613   9 min read         1.5K  18 responses
```

20745 13 min read

298 3 responses

## 1.1 Inspect missing columns

```
[4]: # Define the function which checks missing data and types of data
def missing_data(data):
    total = data.isnull().sum()
    percent = (data.isnull().sum()/data.isnull().count()*100)
    tt = pd.concat([total, percent], axis=1, keys=['Total', 'Percent'])
    types = []
    for col in data.columns:
        dtype = str(data[col].dtype)
        types.append(dtype)
    tt['Types'] = types
    return(np.transpose(tt))
```

```
[5]: missing_data(articles)
```

```
[5]:
```

	author	linkOfAuthorProfile	articleTitle	articleLink	postingTime	\
Total	0	0	1461	0	0	
Percent	0.0	0.0	3.307151	0.0	0.0	
Types	object	object	object	object	object	

  

	minToRead	recommendations	responses
Total	1	268	22235
Percent	0.002264	0.606651	50.331621
Types	object	object	object

## 1.2 There's only one entry that doesn't have minToRead. It turns out to be a navigation article. So I decided to drop it.

```
[6]: articles[articles['minToRead'].isnull()]
```

```
[6]:
```

	author	linkOfAuthorProfile	\
40490	TDS Editors	<a href="https://towardsdatascience.com/@towardsdatasci...">https://towardsdatascience.com/@towardsdatasci...</a>	

  

	articleTitle	articleLink	\
40490	Navigation	<a href="https://towardsdatascience.com/navigation-1f82...">https://towardsdatascience.com/navigation-1f82...</a>	

  

	postingTime	minToRead	recommendations	responses
40490	Nov 14, 2020	None	298	None

```
[7]: articles = articles.dropna(subset=['minToRead'])
```

### 1.3 Inspect articleTitle that has missing values

```
[8]: articles[articles['articleTitle'].isnull()].sample(5)
```

```
[8]:
```

	author \	linkOfAuthorProfile	articleTitle \
28876	Irfan Alghani Khalid	https://towardsdatascience.com/@irfanalghani11...	None
41025	Rose Day	https://towardsdatascience.com/@rjday?source=c...	None
20910	Jo Stichbury	https://towardsdatascience.com/@fluffymaccoy?s...	None
24640	Oleksii Sheremet	https://towardsdatascience.com/@dynamic_phlox_...	None
40737	Sidney Kung	https://towardsdatascience.com/@sidneykung?sou...	None

  

	articleLink	postingTime \
28876	https://towardsdatascience.com/this-is-how-i-w...	May 12, 2020
41025	https://towardsdatascience.com/understanding-t...	Nov 7, 2020
20910	https://towardsdatascience.com/anzograph-a-w3c...	Feb 8, 2019
24640	https://towardsdatascience.com/intersection-ov...	Jul 24, 2020
40737	https://towardsdatascience.com/adapting-data-s...	Nov 19, 2020

  

	minToRead	recommendations	responses
28876	4 min read	294	1 response
41025	6 min read	48	None
20910	7 min read	77	None
24640	3 min read	43	2 responses
40737	6 min read	192	None

#### 1.3.1 I inspected the 14376th entry and I decided to fill the column articleTitle's null entries with segments from articleLink

```
[9]: link = articles.loc[14376].articleLink
link
```

```
[9]: 'https://towardsdatascience.com/data-science-powered-segmentation-models-ae89f9bd405f?source=collection_archive-----6-----'
```

```
[10]: pt1 = link.split("?")[0]
pt1
```

```
[10]: 'https://towardsdatascience.com/data-science-powered-segmentation-models-ae89f9bd405f'
```

```
[11]: pt2 = pt1.split("/")[-1]
      pt2
```

```
[11]: 'data-science-powered-segmentation-models-ae89f9bd405f'
```

```
[12]: pt3 = pt2.split("-")[:-1]
      pt3
```

```
[12]: ['data', 'science', 'powered', 'segmentation', 'models']
```

```
[13]: title = " ".join(pt3)
      title
```

```
[13]: 'data science powered segmentation models'
```

### 1.3.2 Merge the above operations and fill null entries in articleTitle

```
[14]: articles['articleTitle'] = articles['articleTitle'].
      ↪fillna(articles['articleLink'].apply(lambda x: " ".join(x.split("?")[0].
      ↪split("/")[-1].split("-")[:-1])))
```

## 1.4 Add user\_id column with the same technique as above(find segments from linkOfAuthorProfile)

```
[15]: # example link that contains `user_id`
      link = articles.loc[10].linkOfAuthorProfile
      link
```

```
[15]: 'https://towardsdatascience.com/@databaseast?source=collection_archive-----
      3-----'
```

```
[16]: link.split('?')[0].split('@')[-1]
```

```
[16]: 'databaseast'
```

```
[17]: articles['user_id'] = articles['linkOfAuthorProfile'].apply(lambda x: x.split('?
      ↪')[0].split('@')[-1])
```

1.5 `postingTime` 's format is either like “Aug 25, 2018” for dates before 2021 or “Jan 13” for dates after 2021.

```
[18]: # Convert this year's data format in "[Month] [day]" to "[Month] [day] [2021]"
def convert_date(x):
    if ',' not in x:
        x += ', 2021'
    return x

articles['postingTime'] = articles['postingTime'].apply(convert_date)

# Convert the data format in "[Month] [day] [year]" to datetime format
articles['postingTime'] = pd.to_datetime(articles['postingTime'], format='%b
→%d, %Y')
```

1.6 `recommendations` columns have either under 1K(e.g. 221) or  $\geq 1$ K(e.g. 1.3K) or null values

```
[19]: # Fill the null entries in "recommendations" with "0"
articles['recommendations'].fillna('0', inplace=True)
```

```
[20]: # Format "3.4K" to "3400" and also transform from string to integer
def convert_recommendations(x):
    if x[-1] == 'K':
        x = int(float(x[:-1]) * 1000)
    else:
        x = int(x)
    return x

articles['recommendations'] = articles['recommendations'].
→apply(convert_recommendations)
```

1.7 `responses` column has either null values or values' format like “2 responses”

```
[21]: # Fill the null entries in `responses` with "0 response"
articles['responses'].fillna('0 response', inplace=True)
```

```
[22]: # Extract the number and format from string to integer
articles['responses'] = articles['responses'].str.split(' ').str[0].astype(int)
```

## 1.8 minToRead column has format “3 min read”

```
[23]: # Extract the number and format from string to integer
articles['minToRead'] = articles['minToRead'].str.split(' ').str[0].astype(int)
```

## 2 Data Preprocessing for users file

```
[24]: profiles = pd.read_json('raw_data/users/users_07_04_2021.json')
```

```
[25]: len(profiles)
```

```
[25]: 8000
```

### 2.1 There are 11684 unique number of user ids collected from articles file but there are only 8000 users' profiles are collected

```
[26]: len(set(articles.user_id))
```

```
[26]: 11684
```

### 2.2 Select unique user\_id and corresponding author

```
[27]: users = articles[["user_id", "author", "linkOfAuthorProfile"]]
```

```
[28]: users = users.drop_duplicates(subset=["user_id"])
```

```
[29]: len(users)
```

```
[29]: 11684
```

### 2.3 Because of the duplicated names, after merging, there are 100 more wrong entries

```
[30]: df = pd.merge(users, profiles, how="left", left_on="author",
    ↪right_on="user_name")
```

```
[31]: df.sample(5)
```

```
[31]:
```

	user_id	author \
10719	srees1988	Sree
5278	ivana-15022	Ivana Kotorchevikj

224	sethweidman	Seth Weidman
7476	samdenlepcha	Samden Lepcha
702	paulbradshaw	Paul Bradshaw

	linkOfAuthorProfile	user_name \
10719	https://towardsdatascience.com/@srees1988?sour...	NaN
5278	https://towardsdatascience.com/@ivana-15022?so...	NaN
224	https://towardsdatascience.com/@sethweidman?so...	Seth Weidman
7476	https://towardsdatascience.com/@samdenlepcha?s...	NaN
702	https://towardsdatascience.com/@paulbradshaw?s...	NaN

	desc	followers
10719	NaN	NaN
5278	NaN	NaN
224	Became a data scientist to "use math to solve ...	992 Followers
7476	NaN	NaN
702	NaN	NaN

```
[32]: len(df)
```

```
[32]: 11748
```

2.4 I inspected all the duplicated user\_name and deleted the wrong entries(I only showed the first inspection and omitted the output of all the other inspections)

```
[33]: duplicated = profiles[profiles.duplicated(subset=['user_name'])].user_name
      duplicated
```

```
[33]: 607      Gagandeep Singh
      2337      Aditya Sharma
      2944      Abhishek Kumar
      3339      Harshdeep Singh
      3888      Ofer Tirosh
      4006      Gaurav
      4015      Harshit Sharma
      4320      Shen Huang
      4343      An Nguyen
      4726      Salil Jain
      4766      Pranjal Gupta
      4881      Shubham Gupta
      4899      Bruno Santos
      5573      Sahil Gupta
      5932      Phoebe Wong
      5989      Ravi Ranjan
      6252      James
```

```

6562      Abhishek Kumar
7038      Nishant Sinha
7059      Vishal Singh
7329      Manu Sharma
7358      Shekhar Koirala
7411      Nick Jones
7421      Justin
7434      Wendy Wong
7551      Jason Lee
7578      Sue Liu
7725      Christina
7776      Mayank Mishra
7887      Benjamin Peterson
7949      Shikhar Gupta
Name: user_name, dtype: object

```

```
[34]: df[df.author=='Gagandeep Singh']
```

```

[34]:      user_id      author \
3233    gaganmanku96  Gagandeep Singh
3234    gaganmanku96  Gagandeep Singh
4395  singh.gagandeep8  Gagandeep Singh
4396  singh.gagandeep8  Gagandeep Singh

      linkOfAuthorProfile      user_name \
3233  https://towardsdatascience.com/@gaganmanku96?s...  Gagandeep Singh
3234  https://towardsdatascience.com/@gaganmanku96?s...  Gagandeep Singh
4395  https://towardsdatascience.com/@singh.gagandee...  Gagandeep Singh
4396  https://towardsdatascience.com/@singh.gagandee...  Gagandeep Singh

      desc      followers
3233  Data Scientist at Zykr. Geeky - 578 Followers
3234    Big Data Engineer at WooliesX  74 Followers
4395  Data Scientist at Zykr. Geeky - 578 Followers
4396    Big Data Engineer at WooliesX  74 Followers

```

```
[35]: df = df.drop(index=[3234, 4395])
```

```
[36]: # df[df.author=='Aditya Sharma']
```

```
[37]: df = df.drop(index=[2474, 8620])
```

```
[38]: # df[df.author=='Abhishek Kumar']
```

```
[39]: df = df.drop(index=[6811, 6813, 2828, 4363, 4364, 2827, 2828])
```

```
[40]: # df[df.author=='Harshdeep Singh']
```



```
[41]: df = df.drop(index=[1180, 3498])
```

```
[42]: # df[df.author=='Ofar Tirosh']
```

```
[43]: df = df.drop(index=[3507, 7540])
```

## 2.5 During inspections I also found some profiles' description wasn't collected so I filled them manually

```
[44]: df.loc[3506].desc = "CEO and Founder of Tomedes, a professional services_  
    ↳provider to Fortune 500 companies around the world specializing in_  
    ↳localization and translation."
```

```
[45]: # df[df.author=='Gaurav']
```

```
[46]: df = df.drop(index=[256, 3993])
```

```
[47]: df.loc[257].desc = "Editor of TapTechie Publication and Tech@Breno"
```

```
[48]: # df[df.author=='Harshit Sharma']
```

```
[49]: df = df.drop(index=[384, 830])
```

```
[50]: # df[df.author=='Shen Huang']
```

```
[51]: df = df.drop(index=[3779, 5970])
```

```
[52]: # df[df.author=='An Nguyen']
```

```
[53]: df = df.drop(index=[955, 6053])
```

```
[54]: # df[df.author=='Salil Jain']
```

```
[55]: df = df.drop(index=[328, 7938])
```

```
[56]: # df[df.author=='Pranjal Gupta']
```

```
[57]: df = df.drop(index=[6699, 8601])
```

```
[58]: # df[df.author=='Shubham Gupta']
```

```
[59]: df = df.drop(index=[8120, 8820])
```

```
[60]: # df[df.author=='Bruno Santos']
```

```

[61]: df = df.drop(index=[2628, 4611])

[62]: # df[df.author=='Sahil Gupta']

[63]: df = df.drop(index=[3909, 8190])

[64]: # df[df.author=='Phoebe Wong']

[65]: df = df.drop(index=[4299, 5801])

[66]: # df[df.author=='Ravi Ranjan']

[67]: df = df.drop(index=[938, 3237])

[68]: # df[df.author=='James']

[69]: df = df.drop(index=[331, 8613])

[70]: # df[df.author=='Nishant Sinha']

[71]: df = df.drop(index=[213, 2663])

[72]: # df[df.author=='Vishal Singh']

[73]: df = df.drop(index=[1002, 1623])

[74]: df.loc[1001].desc = 'Medium member since August 2020'

[75]: # df[df.author=='Manu Sharma']

[76]: df = df.drop(index=[4068, 5522])

[77]: # df[df.author=='Shekhar Koirala']

[78]: df = df.drop(index=[894, 1652])

[79]: # df[df.author=='Nick Jones']

[80]: df = df.drop(index=[1054, 2063])

[81]: # df[df.author=='Justin']

[82]: df = df.drop(index=[2371, 5415])

[83]: df.loc[2372].desc="Hello, world! My name is Justin. I solve problems using data.
    ↪ Check me out at embracingtherandom.com and linkedin.com/in/justin-m-evans/"

```

```
[84]: # df[df.author=='Wendy Wong']
```

```
[85]: df = df.drop(index=[1519, 4695])
```

```
[86]: # df[df.author=='Jason Lee']
```

```
[87]: df = df.drop(index=[3501, 8042])
```

```
[88]: # df[df.author=='Sue Liu']
```

```
[89]: df = df.drop(index=[1537, 7131])
```

```
[90]: # df[df.author=='Christina']
```

```
[91]: df = df.drop(index=[932, 8556])
```

```
[92]: # df[df.author=='Mayank Mishra']
```

```
[93]: df = df.drop(index=[7293, 7555])
```

```
[94]: # df[df.author=='Benjamin Peterson']
```

```
[95]: df = df.drop(index=[3511, 6446])
```

```
[96]: # df[df.author=='Shikhar Gupta']
```

```
[97]: df = df.drop(index=[47, 6058])
```

**2.6 Now there are no duplicated wrong entries!!! I also dropped the duplicated column user\_name**

```
[98]: df[df.duplicated(subset=["user_id"])]
```

```
[98]: Empty DataFrame
      Columns: [user_id, author, linkOfAuthorProfile, user_name, desc, followers]
      Index: []
```

```
[99]: len(df)
```

```
[99]: 11684
```

```
[100]: df = df.drop(columns="user_name")
```

## 2.7 followers column has null values or that format “552 followers”

```
[101]: # Change the 'null' entries to '0 follower'
df['followers'].fillna('0 follower', inplace=True)

[102]: # Transform the format from "[num] follower(s)" to "num" in integer
df['followers'] = df['followers'].str.split(' ').str[0]

[103]: # Format "3.4K" to "3400" and convert string to integer
def convert_followers(x):
    if x[-1] == 'K':
        x = int(float(x[:-1]) * 1000)
    else:
        x = int(x)
    return x

df['followers'] = df['followers'].apply(convert_followers)

[104]: df.sample(5)
```

```
[104]:
```

	user_id	author \		linkOfAuthorProfile \		desc	followers
9868	jeffrey-scholz	Jeffrey Scholz		https://towardsdatascience.com/@jeffrey-scholz...		NaN	0
6702	dlite	Derek Haynes		https://towardsdatascience.com/@dlite?source=c...		Working on	336
1253	leofle	Lio Fleishman		https://towardsdatascience.com/@leofle?source=...		Partnership Solutions Engineer at Sisense , th...	22
3366	tolaniadekoya	Tolani Adekoya		https://towardsdatascience.com/@tolaniadekoya?...		NaN	0
11435	colefp	Cole		https://towardsdatascience.com/@colefp?source=...		NaN	0

### 3 Export the cleaned data to csv files

```
[105]: df.to_csv("cleaned_data/users/users_07_04_2021.csv", index=False)
```

```
[106]: articles = articles.drop(columns=["author", "linkOfAuthorProfile"])  
articles.to_csv("cleaned_data/articles/articles_07_04_2021.csv", index=False)
```